

Data Intensive Computing - Review Questions 1

Deadline: September 8, 2019

1. Explain how a file region can be left in an inconsistent state in GFS?

2. Briefly explain how HopsFS uses User Defined Partitioning (UDP) and Distributed Aware Transaction (DAT) to improve the system performance (e.g., CPU usage, network traffic and latency)?

3. Show with an example that in the CAP theorem, we can have only two properties out of Consistency, Availability, and Partition Tolerance at the same time.

4. How does BigTable provide strong consistency?

5. Write a simple query in Cypher for each of the following requests:

- Match a Person called “John Doe”
- Find FRIENDS_OF “John Doe”
- Count “John Doe”’s direct acquaintances