



Distributed Machine Learning Frameworks

Amir H. Payberah
payberah@kth.se
2020-12-07





The Course Web Page

<https://fid3024.github.io>

Review of the Current Frameworks



TensorFlow (1/4)

- ▶ TensorFlow supports data parallelism and model partitioning (as of v0.8).



TensorFlow (1/4)

- ▶ TensorFlow supports data parallelism and model partitioning (as of v0.8).
- ▶ As of v2.2, the Multi Worker Mirrored Strategy (allreduce) is integrated into TensorFlow for data parallelism.



TensorFlow (1/4)

- ▶ TensorFlow supports data parallelism and model partitioning (as of v0.8).
- ▶ As of v2.2, the Multi Worker Mirrored Strategy (allreduce) is integrated into TensorFlow for data parallelism.
 - Its update rule is synchronous and it has communication and computation overlapped.



TensorFlow (1/4)

- ▶ TensorFlow supports data parallelism and model partitioning (as of v0.8).
- ▶ As of v2.2, the Multi Worker Mirrored Strategy (allreduce) is integrated into TensorFlow for data parallelism.
 - Its update rule is synchronous and it has communication and computation overlapped.
- ▶ TensorFlow also has extensions to support different parallelization approaches.



TensorFlow (2/4)

- ▶ **Mesh-TensorFlow** is a **language** for distributed deep learning in **TensorFlow**.



TensorFlow (2/4)

- ▶ **Mesh-TensorFlow** is a **language** for distributed deep learning in **TensorFlow**.
- ▶ It is capable of specifying a broad class of **distributed tensor computations**.



TensorFlow (2/4)

- ▶ **Mesh-TensorFlow** is a **language** for distributed deep learning in **TensorFlow**.
- ▶ It is capable of specifying a broad class of **distributed tensor computations**.
- ▶ Mainly used for **model parallelism** in TensorFlow.



TensorFlow (2/4)

- ▶ **Mesh-TensorFlow** is a **language** for distributed deep learning in **TensorFlow**.
- ▶ It is capable of specifying a broad class of **distributed tensor computations**.
- ▶ Mainly used for **model parallelism** in TensorFlow.
- ▶ A **mesh** is an **n-dimensional array** of **processors**, connected by a network.



TensorFlow (2/4)

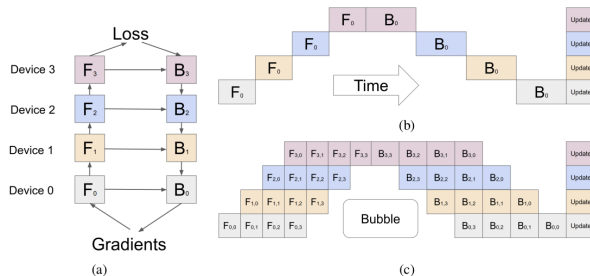
- ▶ **Mesh-TensorFlow** is a **language** for distributed deep learning in **TensorFlow**.
- ▶ It is capable of specifying a broad class of **distributed tensor computations**.
- ▶ Mainly used for **model parallelism** in TensorFlow.
- ▶ A **mesh** is an **n-dimensional array** of **processors**, connected by a network.
- ▶ Each **tensor** is distributed across **all processors in a mesh**.

-
- Figure 1 consists of three sub-diagrams illustrating the proposed framework:
- (a) **Dataflow of the framework**: A vertical stack of four devices (Device 0 to Device 3). Each device has a forward pass (F) and a backward pass (B). The forward pass flows from Device 0 to Device 3, and the backward pass flows from Device 3 to Device 0. A 'Loss' is calculated at Device 3, and 'Gradients' are propagated back from Device 3 to Device 0.
 - (b) **Timeline of the framework**: A sequence of operations (F and B) over time. The operations are arranged in a staircase pattern, showing the progression of the forward and backward passes across the devices.
 - (c) **Timeline of the framework**: A more detailed sequence of operations (F and B) over time, including a 'Bubble' operation. The operations are arranged in a staircase pattern, showing the progression of the forward and backward passes across the devices.

5 / 65

TensorFlow (3/4)

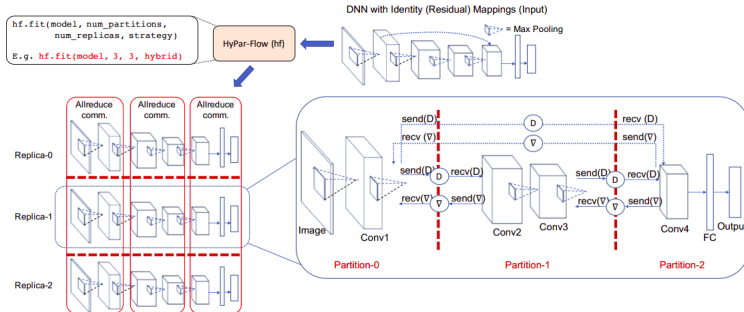
- ▶ **GPipe** is a **pipeline parallelism** library implemented under **Lingvo** (a TensorFlow framework focusing on seq-to-seq models).
- ▶ **Partitions operation** in the **forward and backward pass** and allows data transfer between neighboring partitions.



[Huang et al., GPipe: Efficient Training of Giant Neural Networks using Pipeline Parallelism, 2019]

TensorFlow (4/4)

- **HyPar-Flow** is an implementation of data, model, and hybrid parallelization on **Eager TensorFlow**.



[Awan et al., HyPar-Flow: Exploiting MPI and Keras for Scalable Hybrid-Parallel DNN Training with TensorFlow, 2020]

-

6 / 65



PyTorch (1/4)

- ▶ Caffe is a DL framework that does not support distributed training out-of-the-box.



PyTorch (1/4)

- ▶ Caffe is a DL framework that does not support distributed training out-of-the-box.
- ▶ Many extensions of Caffe to support distributed training centralized or decentralized.



PyTorch (1/4)

- ▶ Caffe is a DL framework that does not support distributed training out-of-the-box.
- ▶ Many extensions of Caffe to support distributed training centralized or decentralized.
- ▶ FireCaffe and MPI-Caffe support data and model parallelism on multi-GPU clusters, respectively.



PyTorch (1/4)

- ▶ Caffe is a DL framework that does not support distributed training out-of-the-box.
- ▶ Many extensions of Caffe to support distributed training centralized or decentralized.
- ▶ FireCaffe and MPI-Caffe support data and model parallelism on multi-GPU clusters, respectively.
- ▶ Intel-Caffe and NUMA-Caffe support data parallelism training on CPU-based clusters.



PyTorch (1/4)

- ▶ **Caffe** is a DL framework that **does not support distributed training** out-of-the-box.
- ▶ Many **extensions** of Caffe to support **distributed training centralized** or **decentralized**.
- ▶ **FireCaffe** and **MPI-Caffe** support **data** and **model parallelism** on **multi-GPU** clusters, respectively.
- ▶ **Intel-Caffe** and **NUMA-Caffe** support **data parallelism** training on **CPU-based** clusters.
- ▶ **S-Caffe** is a **CUDA-Aware MPI runtime** and Caffe for **data parallelism** on **GPU** clusters.



PyTorch (2/4)

- ▶ **Chainer** is a **Define-by-Run (imperative)** DL framework.



PyTorch (2/4)

- ▶ **Chainer** is a **Define-by-Run (imperative)** DL framework.
- ▶ It only supports **data parallelism**.



PyTorch (2/4)

- ▶ **Chainer** is a **Define-by-Run (imperative)** DL framework.
- ▶ It only supports **data parallelism**.
- ▶ It has a **synchronous decentralized** design for **allreduce** communication.



PyTorch (3/4)

- ▶ **PyTorch** is a **successor** of **Caffe2**, which is inspired by **Chainer**.



PyTorch (3/4)

- ▶ **PyTorch** is a **successor** of **Caffe2**, which is inspired by **Chainer**.
- ▶ It is an **imperative** DL framework using **dynamic computation graphs** and **automatic differentiation**.



PyTorch (3/4)

- ▶ **PyTorch** is a **successor** of **Caffe2**, which is inspired by **Chainer**.
- ▶ It is an **imperative** DL framework using **dynamic computation graphs** and **automatic differentiation**.
- ▶ PyTorch mainly focuses on **ease of use**, and enables users with options in training their models.



PyTorch (3/4)

- ▶ **PyTorch** is a **successor** of **Caffe2**, which is inspired by **Chainer**.
- ▶ It is an **imperative** DL framework using **dynamic computation graphs** and **automatic differentiation**.
- ▶ PyTorch mainly focuses on **ease of use**, and enables users with options in training their models.
- ▶ **PyTorch RPC** is developed to support **model parallelism**.



PyTorch (4/4)

- ▶ **PyTorch Distributed Data Parallel (DPP)** is an extra feature to **PyTorch** (available as of v1.5).



PyTorch (4/4)

- ▶ **PyTorch Distributed Data Parallel (DDP)** is an extra feature to **PyTorch** (available as of v1.5).
- ▶ **PyTorch DDP** utilizes some techniques to **increase performance**, such as



PyTorch (4/4)

- ▶ **PyTorch Distributed Data Parallel (DDP)** is an extra feature to **PyTorch** (available as of v1.5).
- ▶ **PyTorch DDP** utilizes some techniques to **increase performance**, such as
 - **Gradient bucketing** (small tensors bucket into one allreduce operation)



PyTorch (4/4)

- ▶ **PyTorch Distributed Data Parallel (DDP)** is an extra feature to **PyTorch** (available as of v1.5).
- ▶ **PyTorch DDP** utilizes some techniques to **increase performance**, such as
 - **Gradient bucketing** (small tensors bucket into one allreduce operation)
 - **Overlapping communication** with **computation**



PyTorch (4/4)

- ▶ **PyTorch Distributed Data Parallel (DDP)** is an extra feature to **PyTorch** (available as of v1.5).
- ▶ **PyTorch DDP** utilizes some techniques to **increase performance**, such as
 - **Gradient bucketing** (small tensors bucket into one allreduce operation)
 - **Overlapping communication** with **computation**
 - **Skipping synchronization**



MXNet (1/2)

- ▶ MXNet is a multi-language ML library.



MXNet (1/2)

- ▶ MXNet is a multi-language ML library.
- ▶ It blends declarative symbolic expression with imperative tensor computation.



MXNet (1/2)


- ▶ **MXNet** is a multi-language ML library.
- ▶ It blends declarative symbolic expression with imperative tensor computation.
- ▶ It uses a distributed key-value store for data synchronization over multiple devices.



MXNet (2/2)

- ▶ **MXNet-MPI** is the extension of MXNet that replaces each worker in a **parameter server architecture** with a **group of workers**.
- ▶ Workers of each group are **synced together** using an **MPI collective** operation.

- ▶ **MXNet-MPI** is the extension of MXNet that replaces each worker in a **parameter server architecture** with a **group of workers**.
- ▶ Workers of each group are **synced together** using an **MPI collective** operation.



[Mamidala et al., MXNet-MPI: Embedding MPI parallelism in Parameter Server Task Model for Scaling Deep Learning, 2018]



Horovod

- ▶ **Horovod** is a **stand-alone** Python library for **data parallelism** using an optimized **ring_allreduce** collective and a **tensor fusion** algorithm.



Horovod

- ▶ **Horovod** is a **stand-alone** Python library for **data parallelism** using an optimized **ring_allreduce** collective and a **tensor fusion** algorithm.
- ▶ It works on top of **another DL framework** (e.g., **TensorFlow**, **PyTorch**, and **MXNET**).



Horovod

- ▶ **Horovod** is a **stand-alone** Python library for **data parallelism** using an optimized **ring_allreduce** collective and a **tensor fusion** algorithm.
- ▶ It works on top of **another DL framework** (e.g., **TensorFlow**, **PyTorch**, and **MXNET**).
- ▶ It has one of the **most optimized asynchronous collectives**.



Horovod


- ▶ **Horovod** is a **stand-alone** Python library for **data parallelism** using an optimized **ring_allreduce** collective and a **tensor fusion** algorithm.
- ▶ It works on top of **another DL framework** (e.g., **TensorFlow**, **PyTorch**, and **MXNET**).
- ▶ It has one of the **most optimized asynchronous collectives**.
- ▶ However, the **communication overhead** significantly grows with the **number of nodes**.



FlexFlow

- FlexFlow can parallelize a DNN in the Sample, Operation, Attribute, and Parameter (SOAP) dimensions.

- FlexFlow can parallelize a DNN in the Sample, Operation, Attribute, and Parameter (SOAP) dimensions.
- It uses guided randomized search of the SOAP space to find a fast parallelization strategy for a specific parallel machine.



[Jia et al., Beyond Data and Model Parallelism for Deep Neural Networks, 2019]



BigDL

- ▶ **BigDL** is a distributed DL framework for **data parallelism** on top of **Spark**.



BigDL

- ▶ **BigDL** is a distributed DL framework for **data parallelism** on top of **Spark**.
- ▶ It does **not support model parallelism**.



BigDL

- ▶ **BigDL** is a distributed DL framework for **data parallelism** on top of **Spark**.
- ▶ It does **not support model parallelism**.
- ▶ It favors **coarse-grained** operations where **data transformations** are **immutable**.



BigDL

- ▶ **BigDL** is a distributed DL framework for **data parallelism** on top of **Spark**.
- ▶ It does **not support model parallelism**.
- ▶ It favors **coarse-grained** operations where **data transformations** are **immutable**.
- ▶ It runs a series of Spark **jobs**, which are **scheduled** by Spark.



BigDL

- ▶ **BigDL** is a distributed DL framework for **data parallelism** on top of **Spark**.
- ▶ It does **not support model parallelism**.
- ▶ It favors **coarse-grained** operations where **data transformations** are **immutable**.
- ▶ It runs a series of Spark **jobs**, which are **scheduled** by Spark.
- ▶ Due to using **Spark**, it is equipped with **fault tolerance** and a **fair load balancing** mechanism.



ZeRO and DeepSpeed

- ▶ **ZeRO** focuses on solving the **memory limitation** problem while attempting to **minimize the overhead**.



ZeRO and DeepSpeed

- ▶ ZeRO focuses on solving the memory limitation problem while attempting to minimize the overhead.
- ▶ It partitions activations, optimizer states, gradients, and parameters and distributes them equally overall available nodes.



ZeRO and DeepSpeed

- ▶ ZeRO focuses on solving the memory limitation problem while attempting to minimize the overhead.
- ▶ It partitions activations, optimizer states, gradients, and parameters and distributes them equally overall available nodes.
- ▶ It then employs overlapping collective operations to reconstruct the tensors as needed.



ZeRO and DeepSpeed

- ▶ **ZeRO** focuses on solving the **memory limitation** problem while attempting to **minimize the overhead**.
- ▶ It partitions **activations**, **optimizer states**, **gradients**, and **parameters** and **distributes them equally** over all available nodes.
- ▶ It then employs **overlapping collective operations** to **reconstruct the tensors** as needed.
- ▶ **DeepSpeed** brings **ZeRO** techniques through lightweight APIs compatible with **PyTorch**.

BigDL: A Distributed Deep Learning Framework for Big Data



Big Data vs. Deep Learning Frameworks

- Big data and deep learning systems have different distributed execution model.



Big Data vs. Deep Learning Frameworks

- ▶ Big data and deep learning systems have different distributed execution model.
- ▶ Big data tasks are embarrassingly parallel and independent of each other.



Big Data vs. Deep Learning Frameworks

- ▶ Big data and deep learning systems have different distributed execution model.
- ▶ Big data tasks are embarrassingly parallel and independent of each other.
- ▶ Deep learning tasks need to coordinate with and depend on others.



Big Data vs. Deep Learning Frameworks

- ▶ Big data and deep learning systems have different distributed execution model.
- ▶ Big data tasks are embarrassingly parallel and independent of each other.
- ▶ Deep learning tasks need to coordinate with and depend on others.
- ▶ Several connectors, e.g., TFX, CaffeOnSpark, TensorFlowOnSpark, SageMaker.




Big Data vs. Deep Learning Frameworks

- ▶ Big data and deep learning systems have different distributed execution model.
- ▶ Big data tasks are embarrassingly parallel and independent of each other.
- ▶ Deep learning tasks need to coordinate with and depend on others.
- ▶ Several connectors, e.g., TFX, CaffeOnSpark, TensorFlowOnSpark, SageMaker.
- ▶ However, the adaptation between different frameworks can impose very large overheads in practice.


Spark Dataflow Model

- **Job** is described based on **directed acyclic graphs (DAG)** **data flow**.




Spark Dataflow Model

- ▶ **Job** is described based on **directed acyclic graphs (DAG)** **data flow**.
- ▶ A **data flow** is composed of any number of **data sources**, **operators**, and **data sinks** by connecting their inputs and outputs.



Spark Dataflow Model

- ▶ **Job** is described based on **directed acyclic graphs (DAG)** **data flow**.
- ▶ A **data flow** is composed of any number of **data sources**, **operators**, and **data sinks** by connecting their inputs and outputs.
- ▶ **Parallelizable operators**



Resilient Distributed Datasets (RDD) (1/2)

- ▶ A distributed memory abstraction.
- ▶ Immutable collections of objects spread across a cluster.
 - Like a `LinkedList <MyObjects>`

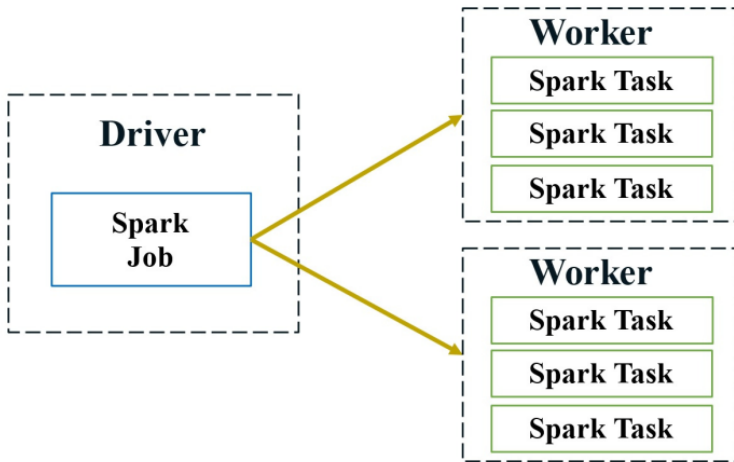


Resilient Distributed Datasets (RDD) (2/2)

- ▶ An **RDD** is divided into a number of **partitions**, which are **atomic** pieces of information.
- ▶ **Partitions** of an RDD can be stored on **different nodes** of a cluster.



Spark Execution Model



[Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019]



BigDL

- ▶ Directly implements the distributed deep learning support in Spark.



BigDL

- ▶ Directly implements the distributed deep learning support in Spark.
- ▶ An data-analytics integrated deep learning pipeline can be executed as a standard Spark jobs.

- ▶ Directly implements the **distributed deep learning** support in **Spark**.
- ▶ An data-analytics **integrated deep learning pipeline** can be executed as a **standard Spark jobs**.

```
#distributed data processing
spark = SparkContext(appName="text_classifier", ...)
input_rdd = spark.textFile("hdfs://...")
train_rdd = input_rdd.map(lambda x: read_text_and_label(x))
                        .map(lambda data: decode_to_ndarrays(data))
                        .map(lambda arrays: to_sample(arrays))

#distributed training
model = Sequential().add(Recurrent().add(LSTM(...)))
                        .add(Linear(...)).add(LogSoftMax())
optimizer = Optimizer(model=model, training_rdd=train_rdd,
                       criterion=ClassNLLCriterion(),
                       optim_method=Adagrad(), ...)
trained_model = optimizer.optimize()

#distributed inference
test_rdd = ...
prediction_rdd = trained_model.predict(test_rdd)
```




Data-Parallel Training in BigDL (1/3)

- ▶ BigDL provides **synchronous data-parallel** training to **train** an NN model.

Data-Parallel Training in BigDL (1/3)


- ▶ BigDL provides **synchronous data-parallel** training to **train** an NN model.
- ▶ **RDD of Samples**, which are automatically **partitioned** across the Spark cluster.



[Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019]

Data-Parallel Training in BigDL (1/3)


- ▶ BigDL provides **synchronous data-parallel** training to **train** an NN model.
- ▶ RDD of **Samples**, which are automatically **partitioned** across the Spark cluster.
- ▶ RDD of **models**, each of which is a **replica** of the original NN model.



[Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019]

Data-Parallel Training in BigDL (2/3)


- In each **iteration**, a single **model forward-backward** Spark job.



[Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019]

Data-Parallel Training in BigDL (2/3)


- ▶ In each **iteration**, a single **model forward-backward** Spark job.
- ▶ Applies the functional **zip operator** to the **co-located partitions** of the two RDDs.



[Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019]

Data-Parallel Training in BigDL (2/3)

- ▶ In each **iteration**, a single **model forward-backward** Spark job.
- ▶ Applies the functional **zip** operator to the **co-located** partitions of the two RDDs.
- ▶ Then, computes the **local gradients** in **parallel** for each **model replica**.



[Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019]

Data-Parallel Training in BigDL (3/3)

Algorithm 1 Data-parallel training in BigDL

```
1: for  $i = 1$  to  $M$  do
2:   //“model forward-backward” job
3:   for each task in the Spark job do
4:     read the latest weights;
5:     get a random batch of data from local Sample partition;
6:     compute local gradients (forward-backward on local model
       replica);
7:   end for
8:   //“parameter synchronization” job
9:   aggregate (sum) all the gradients;
10:  update the weights per specified optimization method;
11: end for
```

[Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019]



Parameter Synchronization in BigDL (1/2)

- ▶ Parameter synchronization based using **parameter server** or **AllReduce** requires **fine-grained** data access.



Parameter Synchronization in BigDL (1/2)

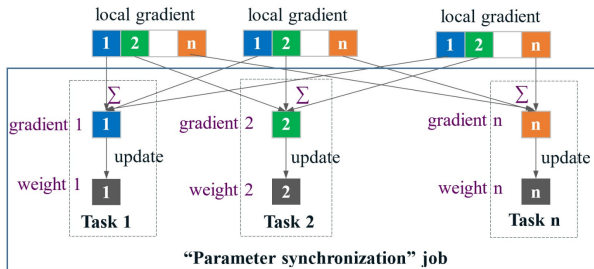
- ▶ Parameter synchronization based using **parameter server** or **AllReduce** requires **fine-grained** data access.
- ▶ **Fine-grained** operations are **not supported** by Spark.



Parameter Synchronization in BigDL (1/2)

- ▶ Parameter synchronization based using parameter server or AllReduce requires fine-grained data access.
- ▶ Fine-grained operations are not supported by Spark.
- ▶ BigDL directly implements an efficient AllReduce-like operation using existing primitives in Spark.

Parameter Synchronization in BigDL (2/2)



Algorithm 2 "Parameter synchronization" job

- 1: **for** each task n in the "parameter synchronization" job **do**
- 2: **shuffle** the n^{th} partition of all gradients to this task;
- 3: aggregate (sum) these gradients;
- 4: updates the n^{th} partition of the weights;
- 5: **broadcast** the n^{th} partition of the updated weights;
- 6: **end for**

[Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019]

PyTorch Distributed: Experiences on Accelerating Data Parallel Training



PyTorch (1/2)

- ▶ **PyTorch** organizes values into **Tensors**, generic **n-dimensional arrays**.



PyTorch (1/2)

- ▶ **PyTorch** organizes values into **Tensors**, generic **n-dimensional arrays**.
- ▶ A **Module** defines a **transform** from input values to output values.



PyTorch (1/2)

- ▶ **PyTorch** organizes values into **Tensors**, generic **n-dimensional arrays**.
- ▶ A **Module** defines a **transform** from input values to output values.
 - In this class, applications provide **their model** at **construction time**.



PyTorch (1/2)

- ▶ **PyTorch** organizes values into **Tensors**, generic **n-dimensional arrays**.
- ▶ A **Module** defines a **transform** from input values to output values.
 - In this class, applications provide **their model** at **construction time**.
 - Its behavior during the **forward pass** is specified by its **forward** member function.



PyTorch (1/2)

- ▶ **PyTorch** organizes values into **Tensors**, generic **n-dimensional arrays**.
- ▶ A **Module** defines a **transform** from input values to output values.
 - In this class, applications provide **their model** at **construction time**.
 - Its behavior during the **forward pass** is specified by its **forward** member function.
- ▶ A **Module** can contain **Tensors** as **parameters**.



PyTorch (1/2)

- ▶ **PyTorch** organizes values into **Tensors**, generic **n-dimensional arrays**.
- ▶ A **Module** defines a **transform** from input values to output values.
 - In this class, applications provide **their model** at **construction time**.
 - Its behavior during the **forward pass** is specified by its **forward** member function.
- ▶ A **Module** can contain **Tensors** as **parameters**.
 - A **LinearModule** contains a **weight** and a **bias** parameter.



PyTorch (1/2)

- ▶ **PyTorch** organizes values into **Tensors**, generic **n-dimensional arrays**.
- ▶ A **Module** defines a **transform** from input values to output values.
 - In this class, applications provide **their model** at **construction time**.
 - Its behavior during the **forward pass** is specified by its **forward** member function.
- ▶ A **Module** can contain **Tensors** as **parameters**.
 - A **LinearModule** contains a **weight** and a **bias** parameter.
 - Whose **forward** function generates the output by **multiplying** the input with the **weight** and **adding** the **bias**.



PyTorch (1/2)

- ▶ **PyTorch** organizes values into **Tensors**, generic **n-dimensional arrays**.
- ▶ A **Module** defines a **transform** from input values to output values.
 - In this class, applications provide **their model** at **construction time**.
 - Its behavior during the **forward pass** is specified by its **forward** member function.
- ▶ A **Module** can contain **Tensors** as **parameters**.
 - A **LinearModule** contains a **weight** and a **bias** parameter.
 - Whose **forward** function generates the output by **multiplying** the input with the **weight** and **adding** the **bias**.
- ▶ An application **composes** its own **Module** by stitching together **Modules** (e.g., linear, convolution) and **Functions** (e.g., relu, pool) in a **forward** function.



PyTorch (2/2)

```
import torch
import torch.nn as nn
import torch.nn.parallel as par
import torch.optim as optim

# initialize torch.distributed properly
# with init_process_group

# setup model and optimizer
net = nn.Linear(10, 10)
opt = optim.SGD(net.parameters(), lr=0.01)

# run forward pass
inp = torch.randn(20, 10)
exp = torch.randn(20, 10)
out = net(inp)

# run backward pass
nn.MSELoss()(out, exp).backward()

# update parameters
opt.step()
```



Data Parallelism in PyTorch (1/4)

- ▶ **PyTorch** provides **distributed data parallel** as an `nn.Module` class.



Data Parallelism in PyTorch (1/4)

- ▶ **PyTorch** provides **distributed data parallel** as an `nn.Module` class.
- ▶ All **replicas** start from the **same initial values** for model parameters.



Data Parallelism in PyTorch (1/4)

- ▶ **PyTorch** provides **distributed data parallel** as an `nn.Module` class.
- ▶ All **replicas** start from the **same initial values** for model parameters.
- ▶ They **synchronize gradients** to keep parameters consistent **across training iterations**.



Data Parallelism in PyTorch (2/4)

- ▶ PyTorch offers **several tools** to facilitate **distributed training**.



Data Parallelism in PyTorch (2/4)

- ▶ PyTorch offers **several tools** to facilitate **distributed training**.
- ▶ `DataParallel` for **data parallel** training on the **same machine**.



Data Parallelism in PyTorch (2/4)

- ▶ PyTorch offers **several tools** to facilitate **distributed training**.
- ▶ `DataParallel` for **data parallel** training on the **same machine**.
- ▶ `DistributedDataParallel` (DDP) for **data parallel** training **across GPUs and machines**.



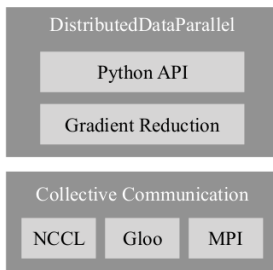
Data Parallelism in PyTorch (2/4)

- ▶ PyTorch offers **several tools** to facilitate **distributed training**.
- ▶ `DataParallel` for **data parallel** training on the **same machine**.
- ▶ `DistributedDataParallel (DDP)` for **data parallel** training **across GPUs and machines**.
- ▶ `RPC` for general distributed **model parallel** training.



Data Parallelism in PyTorch (3/4)

- ▶ DDP module enables data parallel training across multiple processes and machines.




[Li et al., PyTorch Distributed: Experiences on Accelerating Data Parallel Training, 2020]



Data Parallelism in PyTorch (3/4)

- ▶ **DDP** module enables **data parallel** training across multiple processes and machines.
- ▶ **AllReduce** is the primitive communication API used by **DDP**.

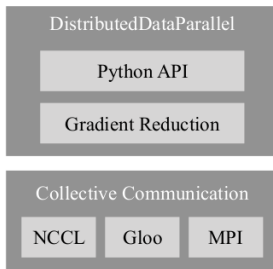


[Li et al., PyTorch Distributed: Experiences on Accelerating Data Parallel Training, 2020]



Data Parallelism in PyTorch (3/4)

- ▶ **DDP** module enables **data parallel** training across multiple processes and machines.
- ▶ **AllReduce** is the primitive communication API used by **DDP**.
- ▶ It is supported by multiple communication libraries, including **NCCL**, **Gloo**, and **MPI**.



[Li et al., PyTorch Distributed: Experiences on Accelerating Data Parallel Training, 2020]



Data Parallelism in PyTorch (4/4)

```
import torch
import torch.nn as nn
import torch.nn.parallel as par
import torch.optim as optim

# initialize torch.distributed properly
# with init_process_group

# setup model and optimizer
net = nn.Linear(10, 10)
opt = optim.SGD(net.parameters(), lr=0.01)

# run forward pass
inp = torch.randn(20, 10)
exp = torch.randn(20, 10)
out = net(inp)

# run backward pass
nn.MSELoss()(out, exp).backward()

# update parameters
opt.step()
```

Data Parallelism in PyTorch (4/4)

```
import torch
import torch.nn as nn
import torch.nn.parallel as par
import torch.optim as optim

# initialize torch.distributed properly
# with init_process_group

# setup model and optimizer
net = nn.Linear(10, 10)
opt = optim.SGD(net.parameters(), lr=0.01)

# run forward pass
inp = torch.randn(20, 10)
exp = torch.randn(20, 10)
out = net(inp)

# run backward pass
nn.MSELoss()(out, exp).backward()

# update parameters
opt.step()
```

```
import torch
import torch.nn as nn
import torch.nn.parallel as par
import torch.optim as optim

# initialize torch.distributed properly
# with init_process_group

# setup model and optimizer
net = nn.Linear(10, 10)
net = par.DistributedDataParallel(net)
opt = optim.SGD(net.parameters(), lr=0.01)

# run forward pass
inp = torch.randn(20, 10)
exp = torch.randn(20, 10)
out = net(inp)

# run backward pass
nn.MSELoss()(out, exp).backward()

# update parameters
opt.step()
```



Gradient Reduction - Naive Solution (1/3)

- ▶ DDP guarantees correctness by letting all training processes:



Gradient Reduction - Naive Solution (1/3)

- ▶ DDP guarantees **correctness** by letting all training processes:
 1. Start from the **same model state**.



Gradient Reduction - Naive Solution (1/3)

- ▶ DDP guarantees **correctness** by letting all training processes:
 1. Start from the **same model state**.
 2. Consume the **same gradients** in every iteration.



Gradient Reduction - Naive Solution (1/3)

- ▶ DDP guarantees **correctness** by letting all training processes:
 1. Start from the **same model state**.
 2. Consume the **same gradients** in every iteration.
- ▶ **Step 1** can be achieved by **broadcasting model states** from one process to all others.



Gradient Reduction - Naive Solution (1/3)

- ▶ DDP guarantees **correctness** by letting all training processes:
 1. Start from the **same model state**.
 2. Consume the **same gradients** in every iteration.
- ▶ **Step 1** can be achieved by **broadcasting model states** from one process to all others.
- ▶ **Step 2** can be achieved by inserting a **gradient synchronization** phase after the **local backward** pass and **before updating local parameters**.



Gradient Reduction - Naive Solution (2/3)

- ▶ To implement the [step 2](#), the PyTorch accepts [custom backward hooks](#).



Gradient Reduction - Naive Solution (2/3)

- ▶ To implement the [step 2](#), the PyTorch accepts [custom backward hooks](#).
- ▶ DDP can register [autograd hooks](#) to [trigger computation](#) after every [backward pass](#).



Gradient Reduction - Naive Solution (2/3)

- ▶ To implement the **step 2**, the PyTorch accepts **custom backward hooks**.
- ▶ DDP can register **autograd hooks** to **trigger computation** after every **backward pass**.
- ▶ When fired, each hook **scans through all local model parameters**, and **retrieves the gradient** tensor from each parameter.



Gradient Reduction - Naive Solution (2/3)

- ▶ To implement the [step 2](#), the PyTorch accepts [custom backward hooks](#).
- ▶ DDP can register [autograd hooks](#) to [trigger computation](#) after every [backward pass](#).
- ▶ When fired, each hook [scans through all local model parameters](#), and [retrieves the gradient](#) tensor from each parameter.
- ▶ Then, it uses the [AllReduce](#) collective communication call to calculate the [average gradients](#) on each parameter across all processes, and writes the result back to the [gradient tensor](#).



Gradient Reduction - Naive Solution (3/3)

- ▶ Two performance concerns:




Gradient Reduction - Naive Solution (3/3)

- ▶ Two performance concerns:
- ▶ 1. Collective communication performs poorly on small tensors, which will be especially prominent on large models with massive numbers of small parameters.


Gradient Reduction - Naive Solution (3/3)

- ▶ Two performance concerns:
 - ▶ 1. Collective communication performs poorly on small tensors, which will be especially prominent on large models with massive numbers of small parameters.
 - ▶ 2. Separating gradient computation and synchronization forfeits the opportunity to overlap computation with communication due to the hard boundary in between.

- Collective communications are more efficient on large tensors.



(a) NCCL



(b) GLOO

[Li et al., PyTorch Distributed: Experiences on Accelerating Data Parallel Training, 2020]



Gradient Reduction - Gradient Bucketing (2/2)

- ▶ Not to launch AllReduce immediately after each gradient tensor becomes available.



Gradient Reduction - Gradient Bucketing (2/2)

- ▶ Not to launch AllReduce immediately after each gradient tensor becomes available.
- ▶ Instead, waits for a short period and buckets multiple gradients into one AllReduce operation.



Gradient Reduction - Gradient Bucketing (2/2)

- ▶ Not to launch `AllReduce` immediately after each gradient tensor becomes available.
- ▶ Instead, waits for a short period and buckets multiple gradients into one `AllReduce` operation.
- ▶ But not to communicate all gradients in one single `AllReduce`, otherwise, no communication can start before the computation is over.

Gradient Reduction - Gradient Bucketing (2/2)

- ▶ Not to launch `AllReduce` immediately after each gradient tensor becomes available.
- ▶ Instead, waits for a short period and buckets multiple gradients into one `AllReduce` operation.
- ▶ But not to communicate all gradients in one single `AllReduce`, otherwise, no communication can start before the computation is over.
- ▶ With relatively small bucket sizes, DDP can launch `AllReduce` operations concurrently with the backward pass to overlap communication with computation.



Overlap Computation with Communication (1/2)

- ▶ AllReduce on gradients can start before the local backward pass finishes.



Overlap Computation with Communication (1/2)

- ▶ `AllReduce` on gradients can **start before** the **local backward pass finishes**.
- ▶ With **bucketing**, `DDP` needs to **wait for all contents** in the **same bucket** before launching communications.



Overlap Computation with Communication (1/2)

- ▶ `AllReduce` on gradients can **start before** the **local backward pass finishes**.
- ▶ With **bucketing**, `DDP` needs to **wait for all contents** in the **same bucket** before launching communications.
- ▶ `DDP` registers one **autograd hook** for **each gradient accumulator**.



Overlap Computation with Communication (1/2)

- ▶ `AllReduce` on gradients can **start before** the **local backward pass finishes**.
- ▶ With **bucketing**, `DDP` needs to **wait for all contents** in the **same bucket** before launching communications.
- ▶ `DDP` registers one **autograd hook** for **each gradient accumulator**.
- ▶ The **hook fires** after its corresponding **accumulator** updating the gradients.




Overlap Computation with Communication (1/2)

- ▶ **AllReduce** on gradients can **start before** the **local backward pass** finishes.
- ▶ With **bucketing**, **DDP** needs to **wait for all contents** in the **same bucket** before launching communications.
- ▶ **DDP** registers one **autograd hook** for **each gradient accumulator**.
- ▶ The **hook fires** after its corresponding **accumulator** updating the gradients.
- ▶ If hooks of **all gradients** in the **same buckets** have fired, then **AllReduce** on that **bucket** will be triggered.

Overlap Computation with Communication (2/2)


- ▶ The **reducing order** must be the **same across** all processes, otherwise, **AllReduce** contents might **mismatch**.



[Li et al., PyTorch Distributed: Experiences on Accelerating Data Parallel Training, 2020]

Overlap Computation with Communication (2/2)


- ▶ The **reducing order** must be the **same across** all processes, otherwise, **AllReduce** contents might **mismatch**.
- ▶ **All processes** must use the **same bucketing order**



[Li et al., PyTorch Distributed: Experiences on Accelerating Data Parallel Training, 2020]

Overlap Computation with Communication (2/2)

- ▶ The **reducing order** must be the **same across** all processes, otherwise, **AllReduce** contents might **mismatch**.
- ▶ **All processes** must use the **same bucketing order**
- ▶ No process can launch **AllReduce** on bucket $i + 1$ **before embarking** bucket i .



[Li et al., PyTorch Distributed: Experiences on Accelerating Data Parallel Training, 2020]



Gradient Accumulation

- ▶ Reduce gradient synchronization frequencies to speed up distributed data parallel training.



Gradient Accumulation

- ▶ Reduce gradient synchronization frequencies to speed up distributed data parallel training.
- ▶ Instead of launching AllReduce in every iteration, it can conduct n local training iterations before synchronizing gradients globally.



Gradient Accumulation

- ▶ Reduce gradient synchronization frequencies to speed up distributed data parallel training.
- ▶ Instead of launching `AllReduce` in every iteration, it can conduct `n` local training iterations before synchronizing gradients globally.
- ▶ Helpful if the input batch is too large to fit into a device.



Gradient Accumulation

- ▶ Reduce gradient synchronization frequencies to speed up distributed data parallel training.
- ▶ Instead of launching `AllReduce` in every iteration, it can conduct `n` local training iterations before synchronizing gradients globally.
- ▶ Helpful if the input batch is too large to fit into a device.
 - It can split one input batch into multiple micro-batches.



Gradient Accumulation

- ▶ Reduce gradient synchronization frequencies to speed up distributed data parallel training.
- ▶ Instead of launching `AllReduce` in every iteration, it can conduct `n` local training iterations before synchronizing gradients globally.
- ▶ Helpful if the input batch is too large to fit into a device.
 - It can split one input batch into multiple micro-batches.
 - Run local forward and backward passes on every micro-batch.

Gradient Accumulation

- ▶ Reduce gradient synchronization frequencies to speed up distributed data parallel training.
- ▶ Instead of launching AllReduce in every iteration, it can conduct n local training iterations before synchronizing gradients globally.
- ▶ Helpful if the input batch is too large to fit into a device.
 - It can split one input batch into multiple micro-batches.
 - Run local forward and backward passes on every micro-batch.
 - Only launch gradient synchronization at the boundaries of large batches.

ZeRO: Memory Optimizations Toward Training Trillion Parameter Models



ZeRO (1/2)

- Data and model parallelisms exhibit fundamental limitations to fit massive models into limited device memory, while obtaining computation, communication and development efficiency.



ZeRO (1/2)

- ▶ Data and model parallelisms exhibit fundamental limitations to fit massive models into limited device memory, while obtaining computation, communication and development efficiency.
- ▶ Zero Redundancy Optimizer (ZeRO) eliminates memory redundancies in data-parallel and model-parallel training.



ZeRO (1/2)

- ▶ Data and model parallelisms exhibit fundamental limitations to fit massive models into limited device memory, while obtaining computation, communication and development efficiency.
- ▶ Zero Redundancy Optimizer (ZeRO) eliminates memory redundancies in data-parallel and model-parallel training.
- ▶ It retains low communication volume and high computational granularity.

ZeRO (1/2)

- ▶ Data and model parallelisms exhibit fundamental limitations to fit massive models into limited device memory, while obtaining computation, communication and development efficiency.
- ▶ Zero Redundancy Optimizer (ZeRO) eliminates memory redundancies in data-parallel and model-parallel training.
- ▶ It retains low communication volume and high computational granularity.
- ▶ Therefore, it allows to scale the model size proportional to the number of devices.



ZeRO (2/2)

- ▶ ZeRO has two sets of optimizations:



ZeRO (2/2)

- ▶ ZeRO has two sets of optimizations:
- ▶ ZeRO-DP (ZeRO Data Parallelism): aims at reducing the memory footprint of the model states.



ZeRO (2/2)

- ▶ ZeRO has two sets of optimizations:
- ▶ ZeRO-DP (ZeRO Data Parallelism): aims at reducing the memory footprint of the model states.
- ▶ ZeRO-R (ZeRO Residual): targetes towards reducing the residual memory consumption.



Where Did All the Memory Go?

- ▶ Model states
 - **Optimizer** states
 - Gradients
 - Parameters
- ▶ **Residual** memory consumption
 - Activations
 - Temporary buffers
 - Memory fragmentation

ZeRO-PD



Optimizing Model State Memory (1/2)

- **Model states** often consume the **largest amount of memory** during **training**.



Optimizing Model State Memory (1/2)

- ▶ **Model states** often consume the **largest amount of memory** during **training**.
 - **Data-Parallel** and **Model-Parallel** approaches **do not** offer satisfying solution.



Optimizing Model State Memory (1/2)

- ▶ **Model states** often consume the **largest amount of memory** during **training**.
 - **Data-Parallel** and **Model-Parallel** approaches **do not** offer satisfying solution.
- ▶ **Data-Parallel (DP)** has **good compute/communication** efficiency, but **poor memory** efficiency.



Optimizing Model State Memory (1/2)

- ▶ **Model states** often consume the **largest amount of memory** during **training**.
 - **Data-Parallel** and **Model-Parallel** approaches **do not** offer satisfying solution.
- ▶ **Data-Parallel (DP)** has **good compute/communication** efficiency, but **poor memory** efficiency.
- ▶ **Model-Parallel (MP)** can have **poor compute/communication** efficiency, but **good memory** efficiency.

Optimizing Model State Memory (1/2)

- ▶ **Model states** often consume the **largest amount of memory** during **training**.
 - **Data-Parallel** and **Model-Parallel** approaches **do not** offer satisfying solution.
- ▶ **Data-Parallel (DP)** has **good compute/communication** efficiency, but **poor memory** efficiency.
- ▶ **Model-Parallel (MP)** can have **poor compute/communication** efficiency, but **good memory** efficiency.
- ▶ Both approaches **maintain all the model states** required over the **entire training process** statically, even though **not all model states** are required all the time during the training.



Optimizing Model State Memory (2/2)

- ▶ ZeRO-DP achieves the computation/communication efficiency of DP, while achieving memory efficiency of MP.



Optimizing Model State Memory (2/2)

- ▶ ZeRO-DP achieves the computation/communication efficiency of DP, while achieving memory efficiency of MP.
- ▶ It removes the memory state redundancies across data-parallel processes by partitioning the model states instead of replicating them.



Optimizing Model State Memory (2/2)

- ▶ ZeRO-DP achieves the computation/communication efficiency of DP, while achieving memory efficiency of MP.
- ▶ It removes the memory state redundancies across data-parallel processes by partitioning the model states instead of replicating them.
- ▶ It retains the compute/communication efficiency by retaining the computational granularity and communication volume of DP using a dynamic communication schedule during training.



Optimization Phases of ZeRO-DP

- ▶ Optimizer state partitioning P_{os}
- ▶ Gradient partitioning P_g
- ▶ Parameter partitioning P_p



Optimizer State Partitioning P_{os}

- ▶ N_d : number of data parallel processes



Optimizer State Partitioning P_{os}

- ▶ N_d : number of data parallel processes
- ▶ Group the optimizer states into N_d equal partitions ($\frac{1}{N_d}$) on each data parallel process.



Optimizer State Partitioning P_{os}

- ▶ N_d : number of data parallel processes
- ▶ Group the optimizer states into N_d equal partitions ($\frac{1}{N_d}$) on each data parallel process.
- ▶ Each data parallel process only updates the its corresponding optimizer states.

Optimizer State Partitioning P_{os}

- ▶ N_d : number of data parallel processes
- ▶ Group the optimizer states into N_d equal partitions ($\frac{1}{N_d}$) on each data parallel process.
- ▶ Each data parallel process only updates the its corresponding optimizer states.
- ▶ Performs an all-gather across the data parallel process at the end of each training step to get the fully updated parameters across all data parallel process.



Gradient Partitioning P_g

- ▶ Each data parallel process only needs the reduced gradients for the corresponding parameters.



Gradient Partitioning P_g

- ▶ Each data parallel process only needs the reduced gradients for the corresponding parameters.
- ▶ As each gradient of each layer becomes available during the backward propagation, only the data parallel process responsible for updating the corresponding parameters will reduce them.



Gradient Partitioning P_g

- ▶ Each data parallel process only needs the reduced gradients for the corresponding parameters.
- ▶ As each gradient of each layer becomes available during the backward propagation, only the data parallel process responsible for updating the corresponding parameters will reduce them.
- ▶ After the reduction, the gradients are no longer needed and their memory can be released.



Gradient Partitioning P_g

- ▶ Each data parallel process only needs the reduced gradients for the corresponding parameters.
- ▶ As each gradient of each layer becomes available during the backward propagation, only the data parallel process responsible for updating the corresponding parameters will reduce them.
- ▶ After the reduction, the gradients are no longer needed and their memory can be released.
- ▶ This is a Reduce-Scatter operation, where gradients corresponding to different parameters are reduced to different process.



Gradient Partitioning P_p

- ▶ Each **process** only stores the **parameters** corresponding to **its partition**.



Gradient Partitioning P_p

- ▶ Each **process** only stores the **parameters** corresponding to **its partition**.
- ▶ When the **parameters outside of its partition** are required for **forward and backward propagation**, they are received from the appropriate data parallel process through **broadcast**.



Gradient Partitioning P_p

- ▶ Each **process** only stores the **parameters** corresponding to **its partition**.
- ▶ When the **parameters outside of its partition** are required for **forward and backward propagation**, they are received from the appropriate data parallel process through **broadcast**.
- ▶ This approach **increases** the **total communication** volume of a baseline DP system to **1.5x**, while enabling **memory reduction** proportional to N_d .

ZeRO-R



Optimizing Residual State Memory

- ▶ ZeRO-DP boosts memory efficiency for **model states**.



Optimizing Residual State Memory

- ▶ ZeRO-DP boosts memory efficiency for **model states**.
- ▶ The **rest of the memory** consumed by **activations**, **temporary buffers**, and unusable **memory fragments**.



Optimizing Residual State Memory

- ▶ ZeRO-DP boosts memory efficiency for **model states**.
- ▶ The **rest of the memory** consumed by **activations**, **temporary buffers**, and unusable **memory fragments**.
- ▶ ZeRO-R optimizes the **residual memory** consumed by the following three factors:



Optimizing Residual State Memory

- ▶ ZeRO-DP boosts memory efficiency for **model states**.
- ▶ The **rest of the memory** consumed by **activations**, **temporary buffers**, and unusable **memory fragments**.
- ▶ ZeRO-R optimizes the **residual memory** consumed by the following three factors:
 1. Optimizes **activation memory** (stored from **forward** pass in order to perform **backward** pass) by **activation partitioning**. It also **offloads** activations to **CPU** when appropriate.



Optimizing Residual State Memory

- ▶ ZeRO-DP boosts memory efficiency for **model states**.
- ▶ The **rest of the memory** consumed by **activations**, **temporary buffers**, and unusable **memory fragments**.
- ▶ ZeRO-R optimizes the **residual memory** consumed by the following three factors:
 1. Optimizes **activation memory** (stored from **forward** pass in order to perform **backward** pass) by **activation partitioning**. It also **offloads** activations to **CPU** when appropriate.
 2. Defines appropriate size for **temporary buffers** to strike for a **balance of memory and computation** efficiency.

Optimizing Residual State Memory

- ▶ ZeRO-DP boosts memory efficiency for **model states**.
- ▶ The **rest of the memory** consumed by **activations**, **temporary buffers**, and unusable **memory fragments**.
- ▶ ZeRO-R optimizes the **residual memory** consumed by the following three factors:
 1. Optimizes **activation memory** (stored from **forward** pass in order to perform **backward** pass) by **activation partitioning**. It also **offloads** activations to **CPU** when appropriate.
 2. Defines appropriate size for **temporary buffers** to strike for a **balance of memory and computation** efficiency.
 3. Proactively manages memory based on the different **lifetime of tensors**, preventing **memory fragmentation**.



Optimization Phases of ZeRO-R

- ▶ Partitioned activation checkpointing P_a
- ▶ Constant size buffers C_B
- ▶ Memory defragmentation M_D



Partitioned Activation Checkpointing P_a

- ▶ **MP** by design requires a replication of the activations.
 - Redundant copies of the activations across model parallel GPUs.



Partitioned Activation Checkpointing P_a

- ▶ **MP** by design requires a **replication** of the activations.
 - **Redundant copies** of the **activations** across **model parallel GPUs**.
- ▶ ZeRO eliminates this redundancy by **partitioning** the activations.

Partitioned Activation Checkpointing P_a

- ▶ MP by design requires a replication of the activations.
 - Redundant copies of the activations across model parallel GPUs.
- ▶ ZeRO eliminates this redundancy by partitioning the activations.
- ▶ Once the forward propagation for a layer of a model is computed, the activations are partitioned across all the model parallel process, until it is needed again during the backpropagation.
- ▶ At this point, ZeRO uses an all-gather operation to re-materialize a replicated copy of the activations.

Partitioned Activation Checkpointing P_a

- ▶ MP by design requires a replication of the activations.
 - Redundant copies of the activations across model parallel GPUs.
- ▶ ZeRO eliminates this redundancy by partitioning the activations.
- ▶ Once the forward propagation for a layer of a model is computed, the activations are partitioned across all the model parallel process, until it is needed again during the backpropagation.
- ▶ At this point, ZeRO uses an all-gather operation to re-materialize a replicated copy of the activations.
- ▶ It works in conjunction with activation checkpointing, storing partitioned activation checkpoints only instead of replicated copies.



Constant Size Buffers C_B

- ▶ ZeRO selects the sizes of the temporal-data buffers to balance memory and compute efficiency.



Constant Size Buffers C_B

- ▶ ZeRO selects the sizes of the temporal-data buffers to balance memory and compute efficiency.
- ▶ During training, the computational efficiency of some operations can be highly dependent on the input size, with larger inputs achieving higher efficiency.

Constant Size Buffers C_B

- ▶ ZeRO selects the sizes of the temporal-data buffers to balance memory and compute efficiency.
- ▶ During training, the computational efficiency of some operations can be highly dependent on the input size, with larger inputs achieving higher efficiency.
- ▶ To get better efficiency, it fuses all the parameters into a single buffer before applying these operations.
- ▶ The memory overhead of the fused buffers is proportional to the model size, and can become inhibiting.

Constant Size Buffers C_B

- ▶ ZeRO selects the sizes of the temporal-data buffers to balance memory and compute efficiency.
- ▶ During training, the computational efficiency of some operations can be highly dependent on the input size, with larger inputs achieving higher efficiency.
- ▶ To get better efficiency, it fuses all the parameters into a single buffer before applying these operations.
- ▶ The memory overhead of the fused buffers is proportional to the model size, and can become inhibiting.
- ▶ To address this issue, ZeRO-R uses a constant-size fused buffer when the model becomes too large.



Memory Defragmentation M_D (1/2)

- ▶ **Memory fragmentation** in model training occurs as a result of **activation checkpointing** and **gradient computation**.



Memory Defragmentation M_D (1/2)

- ▶ Memory fragmentation in model training occurs as a result of activation checkpointing and gradient computation.
- ▶ During the forward propagation with activation checkpointing, only selected activations are stored for back propagation.



Memory Defragmentation M_D (1/2)

- ▶ Memory fragmentation in model training occurs as a result of activation checkpointing and gradient computation.
- ▶ During the forward propagation with activation checkpointing, only selected activations are stored for back propagation.
 - Most activations are discarded as they can be recomputed again during the back propagation.

Memory Defragmentation M_D (1/2)

- ▶ Memory fragmentation in model training occurs as a result of activation checkpointing and gradient computation.
- ▶ During the forward propagation with activation checkpointing, only selected activations are stored for back propagation.
 - Most activations are discarded as they can be recomputed again during the back propagation.
 - This creates an interleaving of short lived memory (discarded activations) and long lived memory (checkpointed activation), leading to memory fragmentation.



Memory Defragmentation M_D (2/2)

- ▶ During the **backward propagation**, the **parameter gradients** are **long lived**, while **activation gradients** and any other buffers required to compute the parameter gradients are **short lived**.



Memory Defragmentation M_D (2/2)

- ▶ During the **backward propagation**, the **parameter gradients** are **long lived**, while **activation gradients** and any other buffers required to compute the parameter gradients are **short lived**.
- ▶ This **interleaving of short term and long term memory** causes memory **fragmentation**.



Memory Defragmentation M_D (2/2)

- ▶ During the **backward propagation**, the **parameter gradients** are **long lived**, while **activation gradients** and any other buffers required to compute the parameter gradients are **short lived**.
- ▶ This **interleaving of short term and long term memory** causes memory **fragmentation**.
- ▶ ZeRO does memory defragmentation on-the-fly by **pre-allocating contiguous memory chunks** for **activation checkpoints and gradients**, and **copying** them over to the **pre-allocated memory** as they are produced.

Summary



Summary

- ▶ BigDL
- ▶ PyTorch Distributed
- ▶ ZeRO

- ▶ Hasheminezhad et al., Towards a Scalable and Distributed Infrastructure for Deep Learning Applications, 2020
- ▶ Dai et al., BigDL: A Distributed Deep Learning Framework for Big Data, 2019
- ▶ Li et al., PyTorch Distributed: Experiences on Accelerating Data Parallel Training, 2020
- ▶ Rajbhandari et al., ZeRO: Memory Optimizations Toward Training Trillion Parameter Models, 2020

Questions?